

dog_app

June 27, 2021

1 Convolutional Neural Networks

1.1 Project: Write an Algorithm for a Dog Identification App

In this notebook, some template code has already been provided for you, and you will need to implement additional functionality to successfully complete this project. You will not need to modify the included code beyond what is requested. Sections that begin with '**(IMPLEMENTATION)**' in the header indicate that the following block of code will require additional functionality which you must provide. Instructions will be provided for each section, and the specifics of the implementation are marked in the code block with a 'TODO' statement. Please be sure to read the instructions carefully!

Note: Once you have completed all of the code implementations, you need to finalize your work by exporting the Jupyter Notebook as an HTML document. Before exporting the notebook to html, all of the code cells need to have been run so that reviewers can see the final implementation and output. You can then export the notebook by using the menu above and navigating to **File -> Download as -> HTML (.html)**. Include the finished document along with this notebook as your submission.

In addition to implementing code, there will be questions that you must answer which relate to the project and your implementation. Each section where you will answer a question is preceded by a '**Question X**' header. Carefully read each question and provide thorough answers in the following text boxes that begin with '**Answer:**'. Your project submission will be evaluated based on your answers to each of the questions and the implementation you provide.

Note: Code and Markdown cells can be executed using the **Shift + Enter** keyboard shortcut. Markdown cells can be edited by double-clicking the cell to enter edit mode.

The rubric contains *optional* "Stand Out Suggestions" for enhancing the project beyond the minimum requirements. If you decide to pursue the "Stand Out Suggestions", you should include the code in this Jupyter notebook.

Step 0: Import Datasets

Make sure that you've downloaded the required human and dog datasets:

Note: if you are using the Udacity workspace, you DO NOT need to re-download these - they can be found in the /data folder as noted in the cell below.

- Download the [dog dataset](#). Unzip the folder and place it in this project's home directory, at the location /dog_images.
- Download the [human dataset](#). Unzip the folder and place it in the home directory, at location /lfw.

Note: If you are using a Windows machine, you are encouraged to use [7zip](#) to extract the folder.

In the code cell below, we save the file paths for both the human (LFW) dataset and dog dataset in the numpy arrays human_files and dog_files.

```
In [1]: import numpy as np
        from glob import glob

        # load filenames for human and dog images
        human_files = np.array(glob("/data/lfw/*/"))
        dog_files = np.array(glob("/data/dog_images/*/"))

        # print number of images in each dataset
        print('There are %d total human images.' % len(human_files))
        print('There are %d total dog images.' % len(dog_files))
```

There are 13233 total human images.

There are 8351 total dog images.

Step 1: Detect Humans

In this section, we use OpenCV's implementation of [Haar feature-based cascade classifiers](#) to detect human faces in images.

OpenCV provides many pre-trained face detectors, stored as XML files on [github](#). We have downloaded one of these detectors and stored it in the haarcascades directory. In the next code cell, we demonstrate how to use this detector to find human faces in a sample image.

```
In [2]: import cv2
        import matplotlib.pyplot as plt
        %matplotlib inline

        # extract pre-trained face detector
        face_cascade = cv2.CascadeClassifier('haarcascades/haarcascade_frontalface_alt.xml')

        # load color (BGR) image
        img = cv2.imread(human_files[0])
        # convert BGR image to grayscale
        gray = cv2.cvtColor(img, cv2.COLOR_BGR2GRAY)

        # find faces in image
        faces = face_cascade.detectMultiScale(gray)

        # print number of faces detected in the image
        print('Number of faces detected:', len(faces))
```

```

# get bounding box for each detected face
for (x,y,w,h) in faces:
    # add bounding box to color image
    cv2.rectangle(img,(x,y),(x+w,y+h),(255,0,0),2)

# convert BGR image to RGB for plotting
cv_rgb = cv2.cvtColor(img, cv2.COLOR_BGR2RGB)

# display the image, along with bounding box
plt.imshow(cv_rgb)
plt.show()

```

Number of faces detected: 1



Before using any of the face detectors, it is standard procedure to convert the images to grayscale. The `detectMultiScale` function executes the classifier stored in `face_cascade` and takes the grayscale image as a parameter.

In the above code, `faces` is a numpy array of detected faces, where each row corresponds to a detected face. Each detected face is a 1D array with four entries that specifies the bounding box of the detected face. The first two entries in the array (extracted in the above code as `x` and `y`) specify the horizontal and vertical positions of the top left corner of the bounding box. The last two entries in the array (extracted here as `w` and `h`) specify the width and height of the box.

1.1.1 Write a Human Face Detector

We can use this procedure to write a function that returns True if a human face is detected in an image and False otherwise. This function, aptly named `face_detector`, takes a string-valued file path to an image as input and appears in the code block below.

```
In [3]: # returns "True" if face is detected in image stored at img_path
def face_detector(img_path):
    img = cv2.imread(img_path)
    gray = cv2.cvtColor(img, cv2.COLOR_BGR2GRAY)
    faces = face_cascade.detectMultiScale(gray)
    return len(faces) > 0
```

1.1.2 (IMPLEMENTATION) Assess the Human Face Detector

Question 1: Use the code cell below to test the performance of the `face_detector` function.

- What percentage of the first 100 images in `human_files` have a detected human face?
- What percentage of the first 100 images in `dog_files` have a detected human face?

Ideally, we would like 100% of human images with a detected face and 0% of dog images with a detected face. You will see that our algorithm falls short of this goal, but still gives acceptable performance. We extract the file paths for the first 100 images from each of the datasets and store them in the numpy arrays `human_files_short` and `dog_files_short`.

Answer: 98% of the human faces were detected in human files but only about 17% of the dog faces were recognized.

```
In [4]: from tqdm import tqdm

human_files_short = human_files[:100]
dog_files_short = dog_files[:100]

##-## Do NOT modify the code above this line. ##-##

## TODO: Test the performance of the face_detector algorithm
## on the images in human_files_short and dog_files_short.
num_human = 0
for img in human_files_short:
    isHuman = face_detector(img)
    if isHuman:
        num_human += 1
percentage = (num_human/len(human_files_short)) * 100
print('Percentage of humans correctly classified as people: {}'.format(percentage))

num_dog = 0
for img in dog_files_short:
    isHuman = face_detector(img)
    if isHuman:
        num_dog += 1
percentage = (num_dog/len(dog_files_short)) * 100
print('Percentage of dogs misclassified as people: {}'.format(percentage))
```

Percentage of humans correctly classified as people: 98.0%
Percentage of dogs misclassified as people: 17.0%

We suggest the face detector from OpenCV as a potential way to detect human images in your algorithm, but you are free to explore other approaches, especially approaches that make use of deep learning :). Please use the code cell below to design and test your own face detection algorithm. If you decide to pursue this *optional* task, report performance on `human_files_short` and `dog_files_short`.

```
In [5]: ### (Optional)
        ### TODO: Test performance of another face detection algorithm.
        ### Feel free to use as many code cells as needed.
```

Step 2: Detect Dogs

In this section, we use a [pre-trained model](#) to detect dogs in images.

1.1.3 Obtain Pre-trained VGG-16 Model

The code cell below downloads the VGG-16 model, along with weights that have been trained on [ImageNet](#), a very large, very popular dataset used for image classification and other vision tasks. ImageNet contains over 10 million URLs, each linking to an image containing an object from one of [1000 categories](#).

```
In [6]: import torch
        import torchvision.models as models

        # define VGG16 model
        VGG16 = models.vgg16(pretrained=True)

        # check if CUDA is available
        use_cuda = torch.cuda.is_available()

        # move model to GPU if CUDA is available
        if use_cuda:
            VGG16 = VGG16.cuda()
```

```
Downloading: "https://download.pytorch.org/models/vgg16-397923af.pth" to /root/.torch/models/vgg16-397923af.pth
100%|| 553433881/553433881 [00:09<00:00, 56234679.37it/s]
```

Given an image, this pre-trained VGG-16 model returns a prediction (derived from the 1000 possible categories in ImageNet) for the object that is contained in the image.

1.1.4 (IMPLEMENTATION) Making Predictions with a Pre-trained Model

In the next code cell, you will write a function that accepts a path to an image (such as 'dogImages/train/001.Affenpinscher/Affenpinscher_00001.jpg') as input and returns the index corresponding to the ImageNet class that is predicted by the pre-trained VGG-16 model. The output should always be an integer between 0 and 999, inclusive.

Before writing the function, make sure that you take the time to learn how to appropriately pre-process tensors for pre-trained models in the [PyTorch documentation](#).

```
In [7]: from PIL import Image
import torchvision.transforms as transforms

def VGG16_predict(img_path):
    """
    Use pre-trained VGG-16 model to obtain index corresponding to
    predicted ImageNet class for image at specified path

    Args:
        img_path: path to an image

    Returns:
        Index corresponding to VGG-16 model's prediction
    """

    ## TODO: Complete the function.
    ## Load and pre-process an image from the given img_path
    ## Return the *index* of the predicted class for that image

    img = Image.open(img_path)

    transform_pipeline = transforms.Compose([transforms.RandomResizedCrop(250),
                                             transforms.ToTensor()])

    img_tensor = transform_pipeline(img)
    img_tensor = img_tensor.unsqueeze(0)

    if torch.cuda.is_available():
        img_tensor = img_tensor.cuda()

    prediction = VGG16(img_tensor)

    if torch.cuda.is_available():
        prediction = prediction.cpu()

    index = prediction.data.numpy().argmax()

    return index
```

1.1.5 (IMPLEMENTATION) Write a Dog Detector

While looking at the [dictionary](#), you will notice that the categories corresponding to dogs appear in an uninterrupted sequence and correspond to dictionary keys 151-268, inclusive, to include all categories from 'Chihuahua' to 'Mexican hairless'. Thus, in order to check to see if an image is predicted to contain a dog by the pre-trained VGG-16 model, we need only check if the pre-trained model predicts an index between 151 and 268 (inclusive).

Use these ideas to complete the `dog_detector` function below, which returns True if a dog is detected in an image (and False if not).

```
In [8]: ### returns "True" if a dog is detected in the image stored at img_path
def dog_detector(img_path):
    ## TODO: Complete the function.
    index = VGG16_predict(img_path)
    if index >= 151 and index <= 268:
        return True
```

1.1.6 (IMPLEMENTATION) Assess the Dog Detector

Question 2: Use the code cell below to test the performance of your `dog_detector` function.

- What percentage of the images in `human_files_short` have a detected dog?
- What percentage of the images in `dog_files_short` have a detected dog?

Answer: 1% of the humans were misclassified whereas 74% of the dogs were correctly classified

```
In [9]: ### TODO: Test the performance of the dog_detector function
### on the images in human_files_short and dog_files_short.
num_human = 0
for img in human_files_short:
    isHuman = dog_detector(img)
    if isHuman:
        num_human += 1
percentage = (num_human/len(human_files_short)) * 100
print('Percentage of humans misclassified as people: {}'.format(percentage))

num_dog = 0
for img in dog_files_short:
    isHuman = dog_detector(img)
    if isHuman:
        num_dog += 1
percentage = (num_dog/len(dog_files_short)) * 100
print('Percentage of correctly classified as people: {}'.format(percentage))
```

Percentage of humans misclassified as people: 1.0%

Percentage of correctly classified as people: 74.0%

We suggest VGG-16 as a potential network to detect dog images in your algorithm, but you are free to explore other pre-trained networks (such as [Inception-v3](#), [ResNet-50](#), etc). Please use

the code cell below to test other pre-trained PyTorch models. If you decide to pursue this *optional* task, report performance on `human_files_short` and `dog_files_short`.

```
In [10]: ### (Optional)
        ### TODO: Report the performance of another pre-trained network.
        ### Feel free to use as many code cells as needed.
```

Step 3: Create a CNN to Classify Dog Breeds (from Scratch)

Now that we have functions for detecting humans and dogs in images, we need a way to predict breed from images. In this step, you will create a CNN that classifies dog breeds. You must create your CNN *from scratch* (so, you can't use transfer learning *yet!*), and you must attain a test accuracy of at least 10%. In Step 4 of this notebook, you will have the opportunity to use transfer learning to create a CNN that attains greatly improved accuracy.

We mention that the task of assigning breed to dogs from images is considered exceptionally challenging. To see why, consider that *even a human* would have trouble distinguishing between a Brittany and a Welsh Springer Spaniel.

Brittany	Welsh Springer Spaniel
----------	------------------------

It is not difficult to find other dog breed pairs with minimal inter-class variation (for instance, Curly-Coated Retrievers and American Water Spaniels).

Curly-Coated Retriever	American Water Spaniel
------------------------	------------------------

Likewise, recall that labradors come in yellow, chocolate, and black. Your vision-based algorithm will have to conquer this high intra-class variation to determine how to classify all of these different shades as the same breed.

Yellow Labrador	Chocolate Labrador
-----------------	--------------------

We also mention that random chance presents an exceptionally low bar: setting aside the fact that the classes are slightly imbalanced, a random guess will provide a correct answer roughly 1 in 133 times, which corresponds to an accuracy of less than 1%.

Remember that the practice is far ahead of the theory in deep learning. Experiment with many different architectures, and trust your intuition. And, of course, have fun!

1.1.7 (IMPLEMENTATION) Specify Data Loaders for the Dog Dataset

Use the code cell below to write three separate [data loaders](#) for the training, validation, and test datasets of dog images (located at `dog_images/train`, `dog_images/valid`, and `dog_images/test`, respectively). You may find [this documentation on custom datasets](#) to be a useful resource. If you

are interested in augmenting your training and/or validation data, check out the wide variety of [transforms](#)!

```
In [11]: import os
         from torchvision import datasets

         ### TODO: Write data loaders for training, validation, and test sets
         ## Specify appropriate transforms, and batch_sizes
         param_transform_resize = 224
         param_transform_crop = 224
         param_data_directory = "/data/dog_images"

         # define transforms for the training data and testing data
         train_transforms = transforms.Compose([transforms.Resize(param_transform_resize),
                                                transforms.CenterCrop(param_transform_crop),
                                                transforms.RandomHorizontalFlip(),
                                                transforms.RandomVerticalFlip(),
                                                transforms.RandomRotation(20),
                                                transforms.ToTensor(),
                                                transforms.Normalize([0.485, 0.456, 0.406],
                                                                      [0.229, 0.224, 0.225])])

         test_transforms = transforms.Compose([transforms.Resize(param_transform_resize),
                                                transforms.CenterCrop(param_transform_crop),
                                                transforms.ToTensor(),
                                                transforms.Normalize([0.485, 0.456, 0.406],
                                                                      [0.229, 0.224, 0.225])])

         # pass transforms in here, then run the next cell to see how the transforms look
         train_data = datasets.ImageFolder( param_data_directory + '/train', transform=train_transforms)
         test_data = datasets.ImageFolder( param_data_directory + '/test', transform=test_transforms)
         valid_data = datasets.ImageFolder( param_data_directory + '/valid', transform=test_transforms)

         trainloader = torch.utils.data.DataLoader( train_data, batch_size=32, shuffle=True, num_workers=4)
         testloader = torch.utils.data.DataLoader( test_data, batch_size=16, shuffle = False, num_workers=4)
         validloader = torch.utils.data.DataLoader( valid_data, batch_size=16, shuffle = False, num_workers=4)

         # create dictionary for all loaders in one
         loaders_scratch = {
             'train': trainloader,
             'valid': validloader,
             'test': testloader
         }
```

Question 3: Describe your chosen procedure for preprocessing the data. - How does your code resize the images (by cropping, stretching, etc)? What size did you pick for the input tensor, and why? - Did you decide to augment the dataset? If so, how (through translations, flips, rotations, etc)? If not, why not?

Answer: Loaded the training, testing and validation datasets and then created data loaders for the same. After this, I resized all image to center cropped, 224 pixel. Randomly adding rotation, horizontal and vertical flips also helps with overfitting the data.

1.1.8 (IMPLEMENTATION) Model Architecture

Create a CNN to classify dog breed. Use the template in the code cell below.

```
In [60]: import torch.nn as nn
import torch.nn.functional as F

# define the CNN architecture
class Net(nn.Module):
    ### TODO: choose an architecture, and complete the class

    def __init__(self):
        super(Net, self).__init__()
        ## Define layers of a CNN

        self.conv1 = nn.Conv2d(3, 16, 3, padding=1)

        self.conv2 = nn.Conv2d(16, 32, 3, padding=1)

        self.conv3 = nn.Conv2d(32, 64, 3, padding=1)

        self.pool = nn.MaxPool2d(2, 2)

        self.fc1 = nn.Linear(28 * 28 * 64, 500)
        self.fc2 = nn.Linear(500, len(train_data.classes))
        self.dropout = nn.Dropout(0.25)
        self.batch_norm = nn.BatchNorm1d(num_features=500)

    def forward(self, x):
        ## Define forward behavior
        x = self.pool(F.relu(self.conv1(x)))
        x = self.pool(F.relu(self.conv2(x)))
        x = self.pool(F.relu(self.conv3(x)))

        x = x.view(x.size(0), -1)
        x = self.dropout(x)
        x = F.relu(self.batch_norm( self.fc1(x)) )
        x = self.dropout(x)
        x = self.fc2(x)
        return x

### You so NOT have to modify the code below this line. ###

# instantiate the CNN
```

```

model_scratch = Net()
print(model_scratch)

# move tensors to GPU if CUDA is available
if use_cuda:
    model_scratch.cuda()

Net(
  (conv1): Conv2d(3, 16, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
  (conv2): Conv2d(16, 32, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
  (conv3): Conv2d(32, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
  (pool): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
  (fc1): Linear(in_features=50176, out_features=500, bias=True)
  (fc2): Linear(in_features=500, out_features=133, bias=True)
  (dropout): Dropout(p=0.25)
  (batch_norm): BatchNorm1d(500, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
)

```

Question 4: Outline the steps you took to get to your final CNN architecture and your reasoning at each step.

Answer:

I have used three convoluted layers, the last layer outputting 133 dog classes and 2D maxpooling. The convoluted layers I have used are `nn.Conv2d(3, 16, 3, padding=1)`, `nn.Conv2d(16, 32, 3, padding=1)`, `nn.Conv2d(32, 64, 3, padding=1)`.

I chose 2D maxpooling to down-sample because it is simple and a choice used often. This will also help in preventing overfitting of features in each layer. For maxpooling, I have gone with `nn.MaxPool2d(2, 2)` and is effective because it will down-sample x and y's dimensions by a factor of 2.

I also add a linear layer at the end to product a 133 dimension output as we need and dropout to prevent overfitting. The dropout layer was implemented to add a bit of randomness to the model since the dropout layer will randomly drop weights into the layer and causes the outputs to be scaled by a factor of $p/(1-p)$ where in this case, we have used a $p = 0.25$. The dropout method has proven to be an effective technique for regularization and preventing the co-adaptation of neurons.

Then finally, the forward function dictates the forward behavior of the model.

1.1.9 (IMPLEMENTATION) Specify Loss Function and Optimizer

Use the next code cell to specify a [loss function](#) and [optimizer](#). Save the chosen loss function as `criterion_scratch`, and the optimizer as `optimizer_scratch` below.

```

In [14]: import torch.optim as optim

        ### TODO: select loss function
        criterion_scratch = nn.CrossEntropyLoss()

        ### TODO: select optimizer

```

```
optimizer_scratch = optim.SGD(model_scratch.parameters(), lr=0.01)

if use_cuda:
    criterion_scratch = criterion_scratch.cuda()
```

1.1.10 (IMPLEMENTATION) Train and Validate the Model

Train and validate your model in the code cell below. [Save the final model parameters](#) at filepath 'model_scratch.pt'.

```
In [16]: def train(n_epochs, loaders, model, optimizer, criterion, use_cuda, save_path):
    """returns trained model"""
    print("start training for {} epochs ...".format(n_epochs))
    # initialize tracker for minimum validation loss
    valid_loss_min = np.Inf

    # exist save-file, load save file
    if os.path.exists(save_path):
        print("load previous saved model ...")
        model.load_state_dict(torch.load(save_path))

    for epoch in range(1, n_epochs+1):
        # initialize variables to monitor training and validation loss
        train_loss = 0.0
        valid_loss = 0.0

        #####
        # train the model #
        #####
        model.train()      # --- set model to train mode
        for batch_idx, (data, target) in enumerate(loaders['train']):
            # move to GPU
            if use_cuda:
                data, target = data.cuda(), target.cuda()
            ## find the loss and update the model parameters accordingly
            ## record the average training loss, using something like
            ## train_loss = train_loss + ((1 / (batch_idx + 1)) * (loss.data - train_loss))
            # -----
            # clear the gradients of all optimized variables
            optimizer.zero_grad()
            # forward pass: compute predicted outputs by passing inputs to the model
            output = model(data)
            # calculate the batch loss
            loss = criterion(output, target)
            # backward pass: compute gradient of the loss with respect to model parameters
            loss.backward()
            # perform a single optimization step (parameter update)
            optimizer.step()
```

```

        # update training loss
        #train_loss = train_loss + ((1 / (batch_idx + 1)) * (loss.data - train_loss))
        train_loss += loss.item()*data.size(0)
        # -----

#####
# validate the model #
#####
model.eval()          # ---- set model to evaluation mode
for batch_idx, (data, target) in enumerate(loaders['valid']):
    # move to GPU
    if use_cuda:
        data, target = data.cuda(), target.cuda()
    ## update the average validation loss
    # -----
    # forward pass: compute predicted outputs by passing inputs to the model
    output = model(data)
    # calculate the batch loss
    loss = criterion(output, target)
    # update average validation loss
    valid_loss += loss.item() * data.size(0)
    # -----

# -----
# calculate average losses
train_loss = train_loss / len(loaders['train'].dataset)
valid_loss = valid_loss / len(loaders['valid'].dataset)
# -----

# print training/validation statistics
print('Epoch: {} \tTraining Loss: {:.6f} \tValidation Loss: {:.6f}'.format( epoch,

## TODO: save the model if validation loss has decreased
# -----
# save model if validation loss has decreased
if valid_loss <= valid_loss_min:
    #print('Validation loss decreased ({:.6f} --> {:.6f}). Saving model ...'.f
    print(' Saving model ...')
    torch.save(model.state_dict(), save_path)
    valid_loss_min = valid_loss
else:
    print("")
# -----

print("done")
# return trained model
return model

```

```

In [17]: # ---Defining Param-----
         from PIL import ImageFile
         ImageFile.LOAD_TRUNCATED_IMAGES = True

         param_epochs = 50

         # train the model
         model_scratch = train(param_epochs, loaders_scratch, model_scratch, optimizer_scratch,
                               criterion_scratch, use_cuda, 'model_scratch.pt')

start training for 50 epochs ...
Epoch: 1      Training Loss: 4.801988      Validation Loss: 4.677920      Saving model ...
Epoch: 2      Training Loss: 4.598128      Validation Loss: 4.563519      Saving model ...
Epoch: 3      Training Loss: 4.469912      Validation Loss: 4.473160      Saving model ...
Epoch: 4      Training Loss: 4.372150      Validation Loss: 4.379837      Saving model ...
Epoch: 5      Training Loss: 4.287302      Validation Loss: 4.327372      Saving model ...
Epoch: 6      Training Loss: 4.215868      Validation Loss: 4.295828      Saving model ...
Epoch: 7      Training Loss: 4.146123      Validation Loss: 4.223012      Saving model ...
Epoch: 8      Training Loss: 4.086580      Validation Loss: 4.209549      Saving model ...
Epoch: 9      Training Loss: 4.018554      Validation Loss: 4.206627      Saving model ...
Epoch: 10     Training Loss: 3.957132      Validation Loss: 4.117373      Saving model ...
Epoch: 11     Training Loss: 3.897040      Validation Loss: 4.100388      Saving model ...
Epoch: 12     Training Loss: 3.846616      Validation Loss: 4.040172      Saving model ...
Epoch: 13     Training Loss: 3.773540      Validation Loss: 4.017182      Saving model ...
Epoch: 14     Training Loss: 3.729786      Validation Loss: 3.966501      Saving model ...
Epoch: 15     Training Loss: 3.657783      Validation Loss: 3.896129      Saving model ...
Epoch: 16     Training Loss: 3.615992      Validation Loss: 3.906564
Epoch: 17     Training Loss: 3.550210      Validation Loss: 3.943464
Epoch: 18     Training Loss: 3.499192      Validation Loss: 3.820334      Saving model ...
Epoch: 19     Training Loss: 3.449738      Validation Loss: 3.879319
Epoch: 20     Training Loss: 3.396656      Validation Loss: 3.826475
Epoch: 21     Training Loss: 3.344506      Validation Loss: 3.836145
Epoch: 22     Training Loss: 3.283201      Validation Loss: 3.776577      Saving model ...
Epoch: 23     Training Loss: 3.231855      Validation Loss: 3.731792      Saving model ...
Epoch: 24     Training Loss: 3.184089      Validation Loss: 3.745505
Epoch: 25     Training Loss: 3.136433      Validation Loss: 3.862862
Epoch: 26     Training Loss: 3.080465      Validation Loss: 3.750353
Epoch: 27     Training Loss: 3.023791      Validation Loss: 3.847730
Epoch: 28     Training Loss: 2.973949      Validation Loss: 3.700025      Saving model ...
Epoch: 29     Training Loss: 2.935419      Validation Loss: 3.648092      Saving model ...
Epoch: 30     Training Loss: 2.872530      Validation Loss: 3.679000
Epoch: 31     Training Loss: 2.818180      Validation Loss: 3.792285
Epoch: 32     Training Loss: 2.746999      Validation Loss: 3.624864      Saving model ...
Epoch: 33     Training Loss: 2.707788      Validation Loss: 3.550603      Saving model ...
Epoch: 34     Training Loss: 2.674686      Validation Loss: 3.615180
Epoch: 35     Training Loss: 2.615340      Validation Loss: 3.544678      Saving model ...
Epoch: 36     Training Loss: 2.555373      Validation Loss: 3.618633
Epoch: 37     Training Loss: 2.520397      Validation Loss: 3.710036

```

Epoch: 38	Training Loss: 2.462733	Validation Loss: 3.517357	Saving model ...
Epoch: 39	Training Loss: 2.417157	Validation Loss: 3.640233	
Epoch: 40	Training Loss: 2.374699	Validation Loss: 3.569021	
Epoch: 41	Training Loss: 2.314197	Validation Loss: 3.624237	
Epoch: 42	Training Loss: 2.256395	Validation Loss: 3.538845	
Epoch: 43	Training Loss: 2.216783	Validation Loss: 3.530317	
Epoch: 44	Training Loss: 2.167819	Validation Loss: 3.545329	
Epoch: 45	Training Loss: 2.102751	Validation Loss: 3.609093	
Epoch: 46	Training Loss: 2.078983	Validation Loss: 3.575150	
Epoch: 47	Training Loss: 2.043333	Validation Loss: 3.529112	
Epoch: 48	Training Loss: 1.985230	Validation Loss: 3.545127	
Epoch: 49	Training Loss: 1.937228	Validation Loss: 3.499937	Saving model ...
Epoch: 50	Training Loss: 1.875049	Validation Loss: 3.565478	

done

1.1.11 (IMPLEMENTATION) Test the Model

Try out your model on the test dataset of dog images. Use the code cell below to calculate and print the test loss and accuracy. Ensure that your test accuracy is greater than 10%.

```
In [18]: def test(loaders, model, criterion, use_cuda):

    # monitor test loss and accuracy
    test_loss = 0.
    correct = 0.
    total = 0.

    model.eval()
    for batch_idx, (data, target) in enumerate(loaders['test']):
        # move to GPU
        if use_cuda:
            data, target = data.cuda(), target.cuda()
        # forward pass: compute predicted outputs by passing inputs to the model
        output = model(data)
        # calculate the loss
        loss = criterion(output, target)
        # update average test loss
        test_loss = test_loss + ((1 / (batch_idx + 1)) * (loss.data - test_loss))
        # convert output probabilities to predicted class
        pred = output.data.max(1, keepdim=True)[1]
        # compare predictions to true label
        correct += np.sum(np.squeeze(pred.eq(target.data.view_as(pred))).cpu().numpy())
        total += data.size(0)

    print('Test Loss: {:.6f}\n'.format(test_loss))

    print('\nTest Accuracy: %2d%% (%2d/%2d)' % (
```

```

100. * correct / total, correct, total))

# call test function
test(loaders_scratch, model_scratch, criterion_scratch, use_cuda)

```

Test Loss: 3.620545

Test Accuracy: 16% (141/836)

Step 4: Create a CNN to Classify Dog Breeds (using Transfer Learning)

You will now use transfer learning to create a CNN that can identify dog breed from images. Your CNN must attain at least 60% accuracy on the test set.

1.1.12 (IMPLEMENTATION) Specify Data Loaders for the Dog Dataset

Use the code cell below to write three separate [data loaders](#) for the training, validation, and test datasets of dog images (located at dogImages/train, dogImages/valid, and dogImages/test, respectively).

If you like, **you are welcome to use the same data loaders from the previous step**, when you created a CNN from scratch.

```

In [19]: ## TODO: Specify data loaders
         loaders_transfer = loaders_scratch.copy()

```

1.1.13 (IMPLEMENTATION) Model Architecture

Use transfer learning to create a CNN to classify dog breed. Use the code cell below, and save your initialized model as the variable `model_transfer`.

```

In [25]: import torchvision.models as models
         import torch.nn as nn

         ## TODO: Specify model architecture
         model_transfer = models.resnet50(pretrained=True)

         for parameter in model_transfer.parameters():
             parameter.requires_grad = False

         model_transfer.fc = nn.Linear(2048, 133, bias=True)

         fc_parameters = model_transfer.fc.parameters()

         for parameter in fc_parameters:
             parameter.requires_grad = True

```



```

    if use_cuda:
        model_transfer = model_transfer.cuda()
    print(model_transfer)

```

```

ResNet(
  (conv1): Conv2d(3, 64, kernel_size=(7, 7), stride=(2, 2), padding=(3, 3), bias=False)
  (bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (relu): ReLU(inplace)
  (maxpool): MaxPool2d(kernel_size=3, stride=2, padding=1, dilation=1, ceil_mode=False)
  (layer1): Sequential(
    (0): Bottleneck(
      (conv1): Conv2d(64, 64, kernel_size=(1, 1), stride=(1, 1), bias=False)
      (bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (conv2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
      (bn2): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (conv3): Conv2d(64, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
      (bn3): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (relu): ReLU(inplace)
      (downsample): Sequential(
        (0): Conv2d(64, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      )
    )
  )
  (1): Bottleneck(
    (conv1): Conv2d(256, 64, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (conv2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
    (bn2): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (conv3): Conv2d(64, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn3): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (relu): ReLU(inplace)
  )
  (2): Bottleneck(
    (conv1): Conv2d(256, 64, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (conv2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
    (bn2): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (conv3): Conv2d(64, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn3): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (relu): ReLU(inplace)
  )
)
  (layer2): Sequential(
    (0): Bottleneck(
      (conv1): Conv2d(256, 128, kernel_size=(1, 1), stride=(1, 1), bias=False)
      (bn1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (conv2): Conv2d(128, 128, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1), bias=False)
      (bn2): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)

```

```

(conv3): Conv2d(128, 512, kernel_size=(1, 1), stride=(1, 1), bias=False)
(bn3): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
(relu): ReLU(inplace)
(downsample): Sequential(
  (0): Conv2d(256, 512, kernel_size=(1, 1), stride=(2, 2), bias=False)
  (1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
)
)
(1): Bottleneck(
  (conv1): Conv2d(512, 128, kernel_size=(1, 1), stride=(1, 1), bias=False)
  (bn1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (conv2): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
  (bn2): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (conv3): Conv2d(128, 512, kernel_size=(1, 1), stride=(1, 1), bias=False)
  (bn3): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (relu): ReLU(inplace)
)
(2): Bottleneck(
  (conv1): Conv2d(512, 128, kernel_size=(1, 1), stride=(1, 1), bias=False)
  (bn1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (conv2): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
  (bn2): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (conv3): Conv2d(128, 512, kernel_size=(1, 1), stride=(1, 1), bias=False)
  (bn3): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (relu): ReLU(inplace)
)
(3): Bottleneck(
  (conv1): Conv2d(512, 128, kernel_size=(1, 1), stride=(1, 1), bias=False)
  (bn1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (conv2): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
  (bn2): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (conv3): Conv2d(128, 512, kernel_size=(1, 1), stride=(1, 1), bias=False)
  (bn3): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (relu): ReLU(inplace)
)
)
(layer3): Sequential(
  (0): Bottleneck(
    (conv1): Conv2d(512, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1), bias=False)
    (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (conv3): Conv2d(256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn3): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (relu): ReLU(inplace)
  )
  (downsample): Sequential(
    (0): Conv2d(512, 1024, kernel_size=(1, 1), stride=(2, 2), bias=False)
    (1): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  )
)

```

```

    )
)
(1): Bottleneck(
  (conv1): Conv2d(1024, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
  (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
  (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (conv3): Conv2d(256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False)
  (bn3): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (relu): ReLU(inplace)
)
(2): Bottleneck(
  (conv1): Conv2d(1024, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
  (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
  (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (conv3): Conv2d(256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False)
  (bn3): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (relu): ReLU(inplace)
)
(3): Bottleneck(
  (conv1): Conv2d(1024, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
  (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
  (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (conv3): Conv2d(256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False)
  (bn3): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (relu): ReLU(inplace)
)
(4): Bottleneck(
  (conv1): Conv2d(1024, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
  (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
  (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (conv3): Conv2d(256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False)
  (bn3): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (relu): ReLU(inplace)
)
(5): Bottleneck(
  (conv1): Conv2d(1024, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
  (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
  (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (conv3): Conv2d(256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False)
  (bn3): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (relu): ReLU(inplace)
)
)
)

```

```

(layer4): Sequential(
  (0): Bottleneck(
    (conv1): Conv2d(1024, 512, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (conv2): Conv2d(512, 512, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1), bias=False)
    (bn2): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (conv3): Conv2d(512, 2048, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn3): BatchNorm2d(2048, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (relu): ReLU(inplace)
    (downsample): Sequential(
      (0): Conv2d(1024, 2048, kernel_size=(1, 1), stride=(2, 2), bias=False)
      (1): BatchNorm2d(2048, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    )
  )
)
(1): Bottleneck(
  (conv1): Conv2d(2048, 512, kernel_size=(1, 1), stride=(1, 1), bias=False)
  (bn1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (conv2): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
  (bn2): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (conv3): Conv2d(512, 2048, kernel_size=(1, 1), stride=(1, 1), bias=False)
  (bn3): BatchNorm2d(2048, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (relu): ReLU(inplace)
)
(2): Bottleneck(
  (conv1): Conv2d(2048, 512, kernel_size=(1, 1), stride=(1, 1), bias=False)
  (bn1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (conv2): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
  (bn2): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (conv3): Conv2d(512, 2048, kernel_size=(1, 1), stride=(1, 1), bias=False)
  (bn3): BatchNorm2d(2048, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (relu): ReLU(inplace)
)
)
(avgpool): AvgPool2d(kernel_size=7, stride=1, padding=0)
(fc): Linear(in_features=2048, out_features=133, bias=True)
)

```

Question 5: Outline the steps you took to get to your final CNN architecture and your reasoning at each step. Describe why you think the architecture is suitable for the current problem.

Answer: ResNet was used as it is a very effective method as it the previous layers are used to feed data into the next layers. This makes the features more powerful and accuracy is improved.

Here I have used transfer learning to train a network that can classify dog images. The classifier part of the model I have used is a Linear model for the fully connected layer with features `nn.Linear(2048, 133, bias=True)`. Seeing as that the fully connected layer is predominantly used for the forward behavior, `nn.Linear` is applicable since it will use 2048 inputs and create 133 outputs for the 133 breeds.

ResNet50 architecture is an excellent model because ResNet network uses a 34-layer plain

network architecture which was developed as a result of VGG-19. Shortcut connections are also added and these shortcut connections then convert the architecture into the residual network.

1.1.14 (IMPLEMENTATION) Specify Loss Function and Optimizer

Use the next code cell to specify a [loss function](#) and [optimizer](#). Save the chosen loss function as `criterion_transfer`, and the optimizer as `optimizer_transfer` below.

```
In [27]: criterion_transfer = nn.CrossEntropyLoss()
         optimizer_transfer = optim.SGD(model_transfer.fc.parameters(), lr=0.001)
```

1.1.15 (IMPLEMENTATION) Train and Validate the Model

Train and validate your model in the code cell below. [Save the final model parameters](#) at filepath `'model_transfer.pt'`.

```
In [28]: # train the model
         model_transfer = train(param_epochs, loaders_transfer, model_transfer, optimizer_transfer,
                                criterion_transfer, use_cuda, 'model_transfer.pt')

         # load the model that got the best validation accuracy (uncomment the line below)
         model_transfer.load_state_dict(torch.load('model_transfer.pt'))
```

start training for 50 epochs ...

Epoch: 1	Training Loss: 4.500599	Validation Loss: 3.817800	Saving model ...
Epoch: 2	Training Loss: 3.735618	Validation Loss: 2.950638	Saving model ...
Epoch: 3	Training Loss: 3.177112	Validation Loss: 2.415431	Saving model ...
Epoch: 4	Training Loss: 2.748735	Validation Loss: 1.981119	Saving model ...
Epoch: 5	Training Loss: 2.437793	Validation Loss: 1.679195	Saving model ...
Epoch: 6	Training Loss: 2.229697	Validation Loss: 1.495493	Saving model ...
Epoch: 7	Training Loss: 2.051385	Validation Loss: 1.313014	Saving model ...
Epoch: 8	Training Loss: 1.900481	Validation Loss: 1.192562	Saving model ...
Epoch: 9	Training Loss: 1.795679	Validation Loss: 1.099309	Saving model ...
Epoch: 10	Training Loss: 1.710554	Validation Loss: 1.037682	Saving model ...
Epoch: 11	Training Loss: 1.624440	Validation Loss: 0.959847	Saving model ...
Epoch: 12	Training Loss: 1.586273	Validation Loss: 0.910881	Saving model ...
Epoch: 13	Training Loss: 1.497395	Validation Loss: 0.878000	Saving model ...
Epoch: 14	Training Loss: 1.444137	Validation Loss: 0.837889	Saving model ...
Epoch: 15	Training Loss: 1.414137	Validation Loss: 0.807969	Saving model ...
Epoch: 16	Training Loss: 1.363771	Validation Loss: 0.793939	Saving model ...
Epoch: 17	Training Loss: 1.352430	Validation Loss: 0.749769	Saving model ...
Epoch: 18	Training Loss: 1.310754	Validation Loss: 0.735718	Saving model ...
Epoch: 19	Training Loss: 1.277903	Validation Loss: 0.731844	Saving model ...
Epoch: 20	Training Loss: 1.246061	Validation Loss: 0.707495	Saving model ...
Epoch: 21	Training Loss: 1.242312	Validation Loss: 0.682573	Saving model ...
Epoch: 22	Training Loss: 1.219294	Validation Loss: 0.672332	Saving model ...
Epoch: 23	Training Loss: 1.190118	Validation Loss: 0.668709	Saving model ...
Epoch: 24	Training Loss: 1.169959	Validation Loss: 0.650566	Saving model ...
Epoch: 25	Training Loss: 1.143648	Validation Loss: 0.654021	

Epoch: 26	Training Loss: 1.125292	Validation Loss: 0.645956	Saving model ...
Epoch: 27	Training Loss: 1.113363	Validation Loss: 0.629227	Saving model ...
Epoch: 28	Training Loss: 1.088679	Validation Loss: 0.622860	Saving model ...
Epoch: 29	Training Loss: 1.084683	Validation Loss: 0.613485	Saving model ...
Epoch: 30	Training Loss: 1.081354	Validation Loss: 0.607068	Saving model ...
Epoch: 31	Training Loss: 1.067878	Validation Loss: 0.612292	
Epoch: 32	Training Loss: 1.035227	Validation Loss: 0.595983	Saving model ...
Epoch: 33	Training Loss: 1.031238	Validation Loss: 0.580314	Saving model ...
Epoch: 34	Training Loss: 1.018086	Validation Loss: 0.584460	
Epoch: 35	Training Loss: 1.024919	Validation Loss: 0.586371	
Epoch: 36	Training Loss: 1.000058	Validation Loss: 0.573042	Saving model ...
Epoch: 37	Training Loss: 0.991055	Validation Loss: 0.574018	
Epoch: 38	Training Loss: 0.975827	Validation Loss: 0.571047	Saving model ...
Epoch: 39	Training Loss: 0.977063	Validation Loss: 0.559331	Saving model ...
Epoch: 40	Training Loss: 0.956751	Validation Loss: 0.561287	
Epoch: 41	Training Loss: 0.955011	Validation Loss: 0.551844	Saving model ...
Epoch: 42	Training Loss: 0.932747	Validation Loss: 0.546677	Saving model ...
Epoch: 43	Training Loss: 0.934123	Validation Loss: 0.540680	Saving model ...
Epoch: 44	Training Loss: 0.930275	Validation Loss: 0.545332	
Epoch: 45	Training Loss: 0.908657	Validation Loss: 0.555367	
Epoch: 46	Training Loss: 0.921612	Validation Loss: 0.536732	Saving model ...
Epoch: 47	Training Loss: 0.898707	Validation Loss: 0.547314	
Epoch: 48	Training Loss: 0.890394	Validation Loss: 0.537725	
Epoch: 49	Training Loss: 0.869518	Validation Loss: 0.535938	Saving model ...
Epoch: 50	Training Loss: 0.890750	Validation Loss: 0.532726	Saving model ...

done

1.1.16 (IMPLEMENTATION) Test the Model

Try out your model on the test dataset of dog images. Use the code cell below to calculate and print the test loss and accuracy. Ensure that your test accuracy is greater than 60%.

```
In [29]: test(loaders_transfer, model_transfer, criterion_transfer, use_cuda)
```

```
Test Loss: 0.549265
```

```
Test Accuracy: 84% (707/836)
```

1.1.17 (IMPLEMENTATION) Predict Dog Breed with the Model

Write a function that takes an image path as input and returns the dog breed (Affenpinscher, Afghan hound, etc) that is predicted by your model.

```
In [55]: ### TODO: Write a function that takes a path to an image as input  

### and returns the dog breed that is predicted by the model.
```

```

# list of class names by index, i.e. a name can be accessed like class_names[0]
class_names = [item[4:].replace("_", " ") for item in train_data.classes]

def predict_breed_transfer(img_path):
    # load the image and return the predicted breed
    image_tensor = image_convert_tensor(img_path)

    if use_cuda:
        image_tensor = image_tensor.cuda()

    y = model_transfer(image_tensor)

    _, preds_tensor = torch.max(y, 1)
    if not use_cuda:
        prediction = np.squeeze(preds_tensor.numpy())
    else:
        prediction = np.squeeze(preds_tensor.cpu().numpy())

    return class_names[prediction]

def print_image(img_path, title="Title"):
    image = Image.open(img_path)
    plt.title(title)
    plt.imshow(image)
    plt.show()

```

Step 5: Write your Algorithm

Write an algorithm that accepts a file path to an image and first determines whether the image contains a human, dog, or neither. Then, - if a **dog** is detected in the image, return the predicted breed. - if a **human** is detected in the image, return the resembling dog breed. - if **neither** is detected in the image, provide output that indicates an error.

You are welcome to write your own functions for detecting humans and dogs in images, but feel free to use the `face_detector` and `human_detector` functions developed above. You are **required** to use your CNN from Step 4 to predict dog breed.

Some sample output for our algorithm is provided below, but feel free to design your own user experience!

1.1.18 (IMPLEMENTATION) Write your Algorithm

In [56]: *### TODO: Write your algorithm.*

Feel free to use as many code cells as needed.

```

def run_app(img_path):
    ## handle cases for a human face, dog, and neither
    if (face_detector(img_path)):
        print("Hello, human!")

```



Sample Human Output

```

predicted_breed = predict_breed_transfer(img_path)
print_image(img_path, title="Predicted breed: {}".format(predicted_breed) )

print("You most closely resemble-")
print(predicted_breed)
elif (dog_detector(img_path)):
    print("Hi, dog!")
    predicted_breed = predict_breed_transfer(img_path)
    print_image(img_path, title="Predicted breed: {}".format(predicted_breed) )

    print("Your breed is-")
    print(predicted_breed)
else:
    print("Oh no! Looks like we weren't able to predict what type of breed you are!")
    print_image(img_path, title="...")
    print("Please try again :)")

In [57]: def image_convert_tensor(image):
    prediction_transforms = transforms.Compose([transforms.Resize(param_transform_resize),
                                                transforms.CenterCrop(param_transform_crop),
                                                transforms.ToTensor(),
                                                transforms.Normalize([0.485, 0.456, 0.406],
                                                                        [0.229, 0.224, 0.225])])

    img_pil = Image.open( image ).convert('RGB')
    img_tensor = prediction_transforms( img_pil )[:3,:,:].unsqueeze(0)

    return img_tensor

def image_convert(tensor):
    """ Display a tensor as an image. """

    image = tensor.to("cpu").clone().detach()
    image = image.numpy().squeeze()
    image = image.transpose(1,2,0)
    image = image * np.array((0.229, 0.224, 0.225)) + np.array((0.485, 0.456, 0.406))

```



```
image = image.clip(0, 1)

return image
```

Step 6: Test Your Algorithm

In this section, you will take your new algorithm for a spin! What kind of dog does the algorithm think that *you* look like? If you have a dog, does it predict your dog's breed accurately? If you have a cat, does it mistakenly think that your cat is a dog?

1.1.19 (IMPLEMENTATION) Test Your Algorithm on Sample Images!

Test your algorithm at least six images on your computer. Feel free to use any images you like. Use at least two human and two dog images.


Question 6: Is the output better than you expected :) ? Or worse :(? Provide at least three possible points of improvement for your algorithm.

Answer: The output is better than I expected. Since the model had a test accuracy of about 84%, I thought the model would make some mistakes but there were none made. One area for improvement could be increasing the number of epochs run, this would also increase the time of training but we could iterate through various number of epochs to find the optimum. Another aspect would be to increase the classes of dogs since this will help in improving model accuracy. Finally, I think using fully connected layers will also help.


```
In [58]: ## TODO: Execute your algorithm from Step 6 on
         ## at least 6 images on your computer.
         ## Feel free to use as many code cells as needed.

         ## suggested code, below
         for file in np.hstack((human_files[:3], dog_files[:3])):
             run_app(file)
```

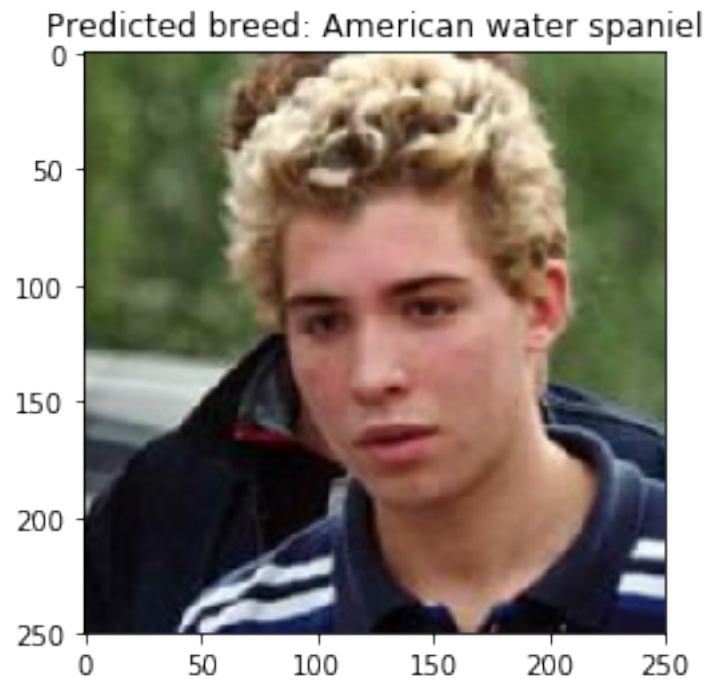
Hello, human!




You most closely resemble-
Chihuahua
Hello, human!




You most closely resemble-
Dachshund
Hello, human!



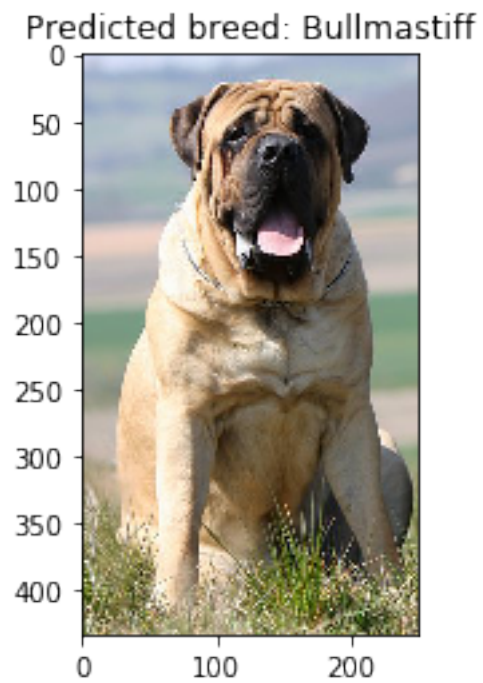
You most closely resemble-
American water spaniel
Hi, dog!



Your breed is-
Bullmastiff
Hi, dog!



```
Your breed is-  
Mastiff  
Hi, dog!
```



```
Your breed is-  
Bullmastiff
```

```
In [ ]:
```